

|               |                             |
|---------------|-----------------------------|
| Student Name  | Sarvesh Naik                |
| SRN No        | 31241943                    |
| Roll No       | 33                          |
| Program       | Computer Engineering        |
| Year          | Third Year                  |
| Division      | F                           |
| Subject       | Artificial Intelligence Lab |
| Assignment No | 7                           |

## Assignment 7

### Problem Statement: Implement the Minimax algorithm to solve the Tic Tac Toe problem

#### Objective:

The objective of this assignment is to implement the Minimax algorithm, an artificial intelligence technique used in decision-making and game playing. The goal is to enable a computer (AI player) to play Tic Tac Toe optimally against an opponent by evaluating all possible future moves and selecting the best one.

This assignment also aims to:

- Understand adversarial search techniques.
- Learn how recursive decision-making works.
- Apply game tree concepts in solving real-world problems.
- Explore optimization using Alpha-Beta Pruning.

#### Theory:

The **Minimax Algorithm** is a decision-making algorithm used in **Artificial Intelligence**, particularly in two-player, turn-based games such as Tic Tac Toe, Chess, and Checkers. It is used to determine the optimal move for a player assuming that the opponent is also playing optimally.

##### 1. Basic Concept

Minimax works on the principle of:

- **Maximizing the player's advantage**
- **Minimizing the opponent's advantage**

There are two types of players:

- **Maximizing Player (MAX)** → tries to maximize the score
- **Minimizing Player (MIN)** → tries to minimize the score

In Tic Tac Toe:

- Player **X** → **Maximizer**
- Player **O** → **Minimizer**

## 2. Game Tree Representation

The algorithm represents the game as a **tree structure**, where:

- Each **node** represents a board state
- Each **edge** represents a possible move
- The **root node** is the current game state
- The **leaf nodes** represent terminal states (win, loss, draw)

The algorithm explores all possible future moves until it reaches terminal states.

## 3. Evaluation Function (Utility Function)

To determine the best move, Minimax uses a scoring system:

- If **MAX (X) wins** → Score = +10
- If **MIN (O) wins** → Score = -10
- If **Draw** → Score = 0

To prefer faster wins and delay losses, depth is considered:

- Score = 10 - depth (for MAX win)
- Score = depth - 10 (for MIN win)

## Working:

The Minimax algorithm works on the principle of **maximizing the AI's chances of winning while minimizing the opponent's chances**.

## Key Concepts:

1. **Players:**
  - **Maximizing Player (X):** AI tries to maximize the score.
  - **Minimizing Player (O):** Opponent tries to minimize the score.
2. **Game Tree:**
  - The algorithm explores all possible game states (moves).
  - Each node represents a board configuration.
3. **Terminal States:**
  - Win → Score assigned (+10 or -10)
  - Draw → Score = 0
4. **Recursive Evaluation:**

- The algorithm explores all possible moves using recursion.
- It assigns scores to each move based on outcomes.
- 5. **Backtracking:**
  - After exploring a move, it backtracks to evaluate other possibilities.
- 6. **Optimal Move Selection:**
  - The best move is chosen based on:
    - Maximum score for AI (X)
    - Minimum score for opponent (O)

## Algorithm:

### Start

#### Initialize the Board

- Represent the Tic Tac Toe board as a list or array of 9 cells.

#### Check Terminal State

- If a player (X or O) has won → return score
- If board is full (draw) → return 0

#### Define Scoring Function

- If **X (Maximizer) wins** → return **+10 – depth**
- If **O (Minimizer) wins** → return **depth – 10**
- If **Draw** → return **0**

#### Recursive Minimax Function

- Input: board, depth, isMaximizing

#### If Maximizing Player (X):

- Set bestScore =  $-\infty$
- For each empty cell:
  - Place 'X' in the cell
  - Recursively call minimax for minimizing player
  - Undo the move (backtrack)
  - Update bestScore = max(bestScore, returned value)
- Return bestScore

#### If Minimizing Player (O):

- Set bestScore =  $+\infty$
- For each empty cell:
  - Place 'O' in the cell
  - Recursively call minimax for maximizing player
  - Undo the move (backtrack)
  - Update bestScore = min(bestScore, returned value)
- Return bestScore

### Find Best Move

- For all possible moves:
  - Apply move for 'X'
  - Call minimax
  - Choose move with highest score

### Apply Alpha-Beta Pruning (Optimization)

- Maintain two values:
  - Alpha ( $\alpha$ ) = best value for maximizer
  - Beta ( $\beta$ ) = best value for minimizer
- While exploring:
  - If  $\beta \leq \alpha$ , stop exploring further (prune branch)

### Return Best Move

End.

### Pseudo Code:

**function MINIMAX(state, depth, isMaximizing):**

**if terminal state:**

**return score of state**

**if isMaximizing:**

**best =  $-\infty$**

**for each move in state:**

**value = MINIMAX(resulting\_state, depth+1, false)**

**best = max(best, value)**

**return best**

**else:**

**best =  $+\infty$**

**for each move in state:**

**value = MINIMAX(resulting\_state, depth+1, true)**

**best = min(best, value)**

**return best**

**function MINIMAX(state, depth, isMaximizing, alpha, beta):**

**if terminal state:**

**return score**

**if isMaximizing:**

**best =  $-\infty$**

**for each move:**

**value = MINIMAX(child, depth+1, false, alpha, beta)**

**best = max(best, value)**

**alpha = max(alpha, best)**

**if  $\beta \leq \alpha$ :**

**break**

**return best**

**else:**

**best =  $+\infty$**

**for each move:**

**value = MINIMAX(child, depth+1, true, alpha, beta)**

**best = min(best, value)**

**beta = min(beta, best)**

**if  $\beta \leq \alpha$ :**

**break**

**return best**

## Code:

```
"""
Minimax Algorithm - Tic Tac Toe

AI Assignment Submission

"""

import matplotlib.pyplot as plt

import matplotlib.gridspec as gridspec

from matplotlib.patches import FancyBboxPatch, FancyArrowPatch, Circle

import matplotlib.path_effects as pe

import matplotlib.patches as mpatches

from matplotlib.lines import Line2D

import numpy as np

import math

import copy

import os

os.makedirs("outputs-7", exist_ok=True)

# -----

# COLOUR PALETTE

# -----

BG = "#0D1117"

PANEL = "#161B22"
```

```
BORDER      = "#30363D"

ACCENT_X    = "#FF6B6B"    # X player (Maximiser) - coral-red

ACCENT_O    = "#4ECDC4"    # O player (Minimiser) - teal

GOLD        = "#FFD166"    # best move highlight

WHITE       = "#E6EDF3"

GREY        = "#8B949E"

GREEN       = "#3FB950"

PURPLE      = "#BC8CFF"

ORANGE      = "#FFA657"
```

```
plt.rcParams.update({

    "font.family": "monospace",

    "text.color":  WHITE,

    "axes.facecolor": BG,

    "figure.facecolor": BG,

})
```

```
# =====
```

```
# CORE MINIMAX ENGINE
```

```
# =====
```

```
def check_winner(board):

    lines = [(0,1,2), (3,4,5), (6,7,8), (0,3,6), (1,4,7), (2,5,8), (0,4,8), (2,4,6)]

    for a,b,c in lines:

        if board[a] == board[b] == board[c] != ' ':
```



```

        return board[a], (a,b,c)

    return None, None

def is_terminal(board):

    winner, _ = check_winner(board)

    return winner is not None or ' ' not in board

def get_available_moves(board):

    return [i for i,v in enumerate(board) if v == ' ']

def minimax(board, depth, is_maximising, alpha=-math.inf, beta=math.inf,
            tree_data=None, parent_id=None, move_made=None):

    """
    Minimax with Alpha-Beta Pruning.

    Optionally records the search tree into `tree_data` (a list of node dicts).

    """

    node_id = len(tree_data) if tree_data is not None else None

    winner, _ = check_winner(board)

    if winner == 'X': score = 10 - depth

    elif winner == 'O': score = depth - 10

    elif ' ' not in board: score = 0

    else: score = None # not terminal

    if tree_data is not None:

```

```

node = {

    "id":          node_id,

    "parent":      parent_id,

    "board":       board[:],

    "depth":       depth,

    "is_max":      is_maximising,

    "score":       score,

    "move":        move_made,

    "pruned":      False,

    "alpha":       alpha,

    "beta":        beta,

}

tree_data.append(node)

if score is not None:

    return score

moves = get_available_moves(board)

best_score = -math.inf if is_maximising else math.inf

best_move = None

for mv in moves:

    new_board = board[:]

    new_board[mv] = 'X' if is_maximising else 'O'

    s = minimax(new_board, depth + 1, not is_maximising,

```

```

        alpha, beta, tree_data, node_id, mv)

    if is_maximising:

        if s > best_score:

            best_score, best_move = s, mv

            alpha = max(alpha, s)

    else:

        if s < best_score:

            best_score, best_move = s, mv

            beta = min(beta, s)

    if beta <= alpha:

        if tree_data is not None:

            tree_data[node_id]["pruned"] = True

        break

    if tree_data is not None:

        tree_data[node_id]["score"] = best_score

        tree_data[node_id]["best_move"] = best_move

    return best_score

def best_move(board):

    """Return the best move index for 'X' (maximiser)."""

    best_sc = -math.inf

```

```

bm = None

for mv in get_available_moves(board):

    b = board[:]

    b[mv] = 'X'

    s = minimax(b, 0, False)

    if s > best_sc:

        best_sc, bm = s, mv

return bm, best_sc

# =====

# HELPER - draw a single board cell grid

# =====

def draw_board_on_ax(ax, board, highlight_cells=None, win_line=None,

                    title="", title_color=WHITE, fontsize=7, cell_size=1):

    ax.set_xlim(0, 3*cell_size); ax.set_ylim(0, 3*cell_size)

    ax.set_aspect('equal')

    ax.axis('off')

    if title:

        ax.set_title(title, color=title_color, fontsize=fontsize,

                    fontweight='bold', pad=2)

    hc = set(highlight_cells) if highlight_cells else set()

    wl = set(win_line) if win_line else set()

```

```

for idx in range(9):

    row, col = divmod(idx, 3)

    y = (2 - row) * cell_size

    x = col * cell_size

    fc = "#1C2733" if idx in hc else ("#0D2818" if idx in wl else PANEL)

    ec = GOLD if idx in hc else (GREEN if idx in wl else BORDER)

    lw = 1.8 if (idx in hc or idx in wl) else 0.6

    rect = FancyBboxPatch((x+0.04*cell_size, y+0.04*cell_size),

                           0.92*cell_size, 0.92*cell_size,

                           boxstyle="round,pad=0.02",

                           facecolor=fc, edgecolor=ec, linewidth=lw)

    ax.add_patch(rect)

    val = board[idx]

    if val != ' ':

        clr = ACCENT_X if val == 'X' else ACCENT_O

        ax.text(x + 0.5*cell_size, y + 0.5*cell_size, val,

                ha='center', va='center',

                fontsize=fontsize*1.8, color=clr,

                fontweight='black',

                path_effects=[pe.withStroke(linewidth=1, foreground=BG)])

```

```

# =====

# FIGURE 1 - ALGORITHM OVERVIEW (2-page explainer)

# =====

def fig_overview():

    fig = plt.figure(figsize=(18, 10), facecolor=BG)

    fig.suptitle("MINIMAX ALGORITHM - Tic Tac Toe",

                 fontsize=22, fontweight='black', color=WHITE, y=0.97)

    gs = gridspec.GridSpec(2, 4, figure=fig, hspace=0.55, wspace=0.35,

                           left=0.04, right=0.96, top=0.90, bottom=0.05)

    # — Left column: pseudocode —————

    ax_code = fig.add_subplot(gs[:, 0])

    ax_code.set_facecolor(PANEL)

    ax_code.axis('off')

    ax_code.set_title("Pseudocode", color=GOLD, fontsize=11, fontweight='bold')

    code = (

        "function MINIMAX(state, depth,\n"

        "                    isMaximising):\n\n"

        "    if TERMINAL(state):\n"

        "        return SCORE(state, depth)\n\n"

        "    if isMaximising:\n"

```

```

"    best ←  $-\infty$ \n"

"    for move in MOVES(state):\n"

"        val ← MINIMAX(APPLY(state,move),\n"

"                        depth+1, FALSE)\n"

"        best ← MAX(best, val)\n"

"    return best\n\n"

" else: # minimising\n"

"    best ←  $+\infty$ \n"

"    for move in MOVES(state):\n"

"        val ← MINIMAX(APPLY(state,move),\n"

"                        depth+1, TRUE)\n"

"        best ← MIN(best, val)\n"

"    return best"

)

ax_code.text(0.05, 0.95, code, transform=ax_code.transAxes,

             fontsize=8.5, color=WHITE, va='top', fontfamily='monospace',

             linespacing=1.55)

# colour-coded legend for pseudocode

for y_, txt, clr in [(0.06, "■ Maximiser (X)", ACCENT_X),

                    (0.03, "■ Minimiser (O)", ACCENT_O)]:

    ax_code.text(0.05, y_, txt, transform=ax_code.transAxes,

                 fontsize=8, color=clr)

# — Scoring rules —————

```

```

ax_score = fig.add_subplot(gs[0, 1])

ax_score.set_facecolor(PANEL); ax_score.axis('off')

ax_score.set_title("Scoring Function", color=GOLD, fontsize=11, fontweight='bold')

rules = [

    ("X wins",    "+10 - depth", ACCENT_X),

    ("O wins",    "depth - 10",  ACCENT_O),

    ("Draw",      "0",           GREY),

    ("X plays",   "Maximise ↑",  ACCENT_X),

    ("O plays",   "Minimise ↓",  ACCENT_O),

]

for i, (lbl, val, clr) in enumerate(rules):

    y_pos = 0.82 - i*0.16

    ax_score.add_patch(FancyBboxPatch((0.04, y_pos-0.06), 0.92, 0.12,

                                     boxstyle="round,pad=0.01",

                                     facecolor=BG, edgecolor=clr, linewidth=1.2,

                                     transform=ax_score.transAxes))

    ax_score.text(0.10, y_pos, lbl, transform=ax_score.transAxes,

                 fontsize=9, color=clr, va='center', fontweight='bold')

    ax_score.text(0.90, y_pos, val, transform=ax_score.transAxes,

                 fontsize=9, color=WHITE, va='center', ha='right')

# — Alpha-Beta box —————

ax_ab = fig.add_subplot(gs[1, 1])

ax_ab.set_facecolor(PANEL); ax_ab.axis('off')

ax_ab.set_title("Alpha-Beta Pruning", color=GOLD, fontsize=11, fontweight='bold')

```



```

ab_text = (

    "Improvement over plain Minimax.\n\n"

    " $\alpha$  = best score Maximiser can guarantee\n"

    " $\beta$  = best score Minimiser can guarantee\n\n"

    "Prune when  $\beta \leq \alpha$ \n"

    "(remaining branches can't affect result)\n\n"

    "Complexity:\n"

    "  Plain Minimax:  $O(b^m)$ \n"

    "  Alpha-Beta:     $O(b^{(m/2)})$  (best case)\n\n"

    "For Tic-Tac-Toe:\n"

    "   $b \approx 5, m = 9 \rightarrow 255,168$  positions\n"

    "  With pruning: significantly fewer nodes"

)

ax_ab.text(0.05, 0.92, ab_text, transform=ax_ab.transAxes,

           fontsize=8.5, color=WHITE, va='top', linespacing=1.5)

# — Example position: X to move —————

example_board = ['X', ' ', 'O',

                 ' ', 'X', ' ',

                 'O', ' ', ' ']

bm_idx, bm_sc = best_move(example_board)

ax_ex = fig.add_subplot(gs[0, 2])

draw_board_on_ax(ax_ex, example_board,

                 highlight_cells=[bm_idx],

```

```

        title=f"Example: X to move → best={bm_sc}",

        title_color=GOLD, fontsize=9)

# — After best move —————

after = example_board[:]

after[bm_idx] = 'X'

winner, wl = check_winner(after)

ax_af = fig.add_subplot(gs[1, 2])

draw_board_on_ax(ax_af, after,

                win_line=list(wl) if wl else None,

                title="After best move → X wins!" if winner == 'X' else "After
best move",

                title_color=GREEN if winner else WHITE, fontsize=9)

# — Game-tree summary —————

ax_info = fig.add_subplot(gs[:, 3])

ax_info.set_facecolor(PANEL); ax_info.axis('off')

ax_info.set_title("Game Tree Facts", color=GOLD, fontsize=11, fontweight='bold')

facts = [

    ("Board states",      "~5,478"),

    ("Terminal nodes",    "255,168"),

    ("X wins",            "131,184"),

    ("O wins",            "77,904"),

    ("Draws",             "46,080"),

```

```

        ("Tree depth",      "max 9"),

        ("Branching factor", "≤ 9"),

        ("Result (optimal)", "Always draw"),

    ]

    for i, (k,v) in enumerate(facts):

        yp = 0.92 - i*0.105

        ax_info.text(0.05, yp, k, transform=ax_info.transAxes,

                     fontsize=9, color=GREY)

        ax_info.text(0.95, yp, v, transform=ax_info.transAxes,

                     fontsize=9, color=WHITE, ha='right', fontweight='bold')

        if i < len(facts)-1:

            line = Line2D([0.04, 0.96], [yp - 0.045, yp - 0.045],

                          color=BORDER, linewidth=0.5,

                          transform=ax_info.transAxes)

            ax_info.add_line(line)

    plt.savefig("outputs-7/fig1_overview.png",

               dpi=150, bbox_inches='tight', facecolor=BG)

    plt.close()

    print("✓ fig1_overview.png")

# =====
# FIGURE 2 - MINIMAX SEARCH TREE (depth-limited, from a mid-game state)
# =====

```

```

def build_tree_limited(board, max_depth=3):

    """Build tree data up to max_depth levels."""

    tree = []

    minimax(board, 0, True, -math.inf, math.inf, tree, None, None)

    return [n for n in tree if n["depth"] <= max_depth]


def fig_search_tree():

    # A near-terminal board so the tree stays manageable

    board = ['X', 'O', 'X',

              'O', 'X', ' ',

              ' ', ' ', 'O']

    tree = []

    minimax(board, 0, True, -math.inf, math.inf, tree, None, None)

    # Layout: assign x positions per depth level

    by_depth = {}

    for n in tree:

        by_depth.setdefault(n["depth"], []).append(n)

    # BFS-based x assignment

    pos = {}

    for d, nodes in sorted(by_depth.items()):

        n_nodes = len(nodes)

        for i, nd in enumerate(nodes):

```

```

        x = (i + 0.5) / n_nodes

        pos[nd["id"]] = (x, 1.0 - d * 0.18)

fig, ax = plt.subplots(figsize=(20, 11), facecolor=BG)

ax.set_facecolor(BG); ax.axis('off')

fig.suptitle("MINIMAX SEARCH TREE (Alpha-Beta Pruning) - X to Move",

             fontsize=16, fontweight='black', color=WHITE, y=0.98)

# Draw edges first

id_map = {n["id"]: n for n in tree}

for nd in tree:

    if nd["parent"] is not None and nd["parent"] in pos:

        x0,y0 = pos[nd["parent"]]

        x1,y1 = pos[nd["id"]]

        clr = ORANGE if nd.get("pruned") else BORDER

        ax.annotate("", xy=(x1,y1), xytext=(x0,y0),

                    arrowprops=dict(arrowstyle="-|>",

                                    color=clr, lw=0.8,

                                    mutation_scale=8))

# Draw nodes

node_w, node_h = 0.055, 0.13

for nd in tree:

    if nd["id"] not in pos:

        continue

```

```

cx, cy = pos[nd["id"]]

# node box colour

if nd.get("pruned"):

    ec = ORANGE; fc = "#2A1A00"

elif nd["is_max"]:

    ec = ACCENT_X; fc = "#2A0A0A"

else:

    ec = ACCENT_O; fc = "#0A2A28"

rect = FancyBboxPatch((cx - node_w/2, cy - node_h/2),

                      node_w, node_h,

                      boxstyle="round,pad=0.005",

                      facecolor=fc, edgecolor=ec,

                      linewidth=1.2,

                      transform=ax.transData)

ax.add_patch(rect)

# tiny board inside node

b = nd["board"]

symbols = {' ': '', 'X': 'X', 'O': 'O'}

colors = {'X': ACCENT_X, 'O': ACCENT_O, ' ': GREY}

cell_w = node_w / 3; cell_h = node_h / 3

for idx in range(9):

    r, c2 = divmod(idx, 3)

```

```

tx = cx - node_w/2 + (c2 + 0.5) * cell_w

ty = cy - node_h/2 + (2 - r + 0.5) * cell_h

# highlight last move

if nd.get("move") == idx:

    ax.add_patch(Circle((tx, ty), cell_w*0.38,

                        facecolor="#333300", edgecolor=GOLD,

                        linewidth=0.5, transform=ax.transData,

                        zorder=3))

if b[idx] != ' ':

    ax.text(tx, ty, b[idx],

            ha='center', va='center', fontsize=5.5,

            color=colors[b[idx]], fontweight='black',

            transform=ax.transData, zorder=4)

# score label

sc = nd.get("score")

if sc is not None:

    sc_clr = GREEN if sc > 0 else (ACCENT_O if sc < 0 else GREY)

    ax.text(cx, cy - node_h/2 - 0.025, f"{sc:+d}" if sc != 0 else "0",

            ha='center', va='top', fontsize=7.5, color=sc_clr,

            fontweight='bold', transform=ax.transData)

# MAX/MIN label on left side

role_txt = "MAX" if nd["is_max"] else "MIN"

role_clr = ACCENT_X if nd["is_max"] else ACCENT_O

```

```

if nd["depth"] == 0:

    ax.text(cx - node_w/2 - 0.01, cy,

            role_txt, ha='right', va='center',

            fontsize=7, color=role_clr, fontweight='bold',

            transform=ax.transData)

# Depth labels on y-axis

for d, nodes in by_depth.items():

    y = 1.0 - d * 0.18

    if nodes:

        role = "MAX (X)" if nodes[0]["is_max"] else "MIN (O)"

        clr = ACCENT_X if nodes[0]["is_max"] else ACCENT_O

        ax.text(0.001, y, f"d={d} {role}",

                ha='left', va='center', fontsize=8,

                color=clr, transform=ax.transAxes)

# Legend

legend_elements = [

    mpatches.Patch(facecolor="#2A0A0A", edgecolor=ACCENT_X, label="MAX node (X  
maximises)"),

    mpatches.Patch(facecolor="#0A2A28", edgecolor=ACCENT_O, label="MIN node (O  
minimises)"),

    mpatches.Patch(facecolor="#2A1A00", edgecolor=ORANGE, label="Pruned branch"),

    Line2D([0],[0], color=GREEN, lw=2, label="Score > 0 (X winning)"),

    Line2D([0],[0], color=ACCENT_O, lw=2, label="Score < 0 (O winning)"),

    Line2D([0],[0], color=GREY, lw=2, label="Score = 0 (Draw)"),

```



```

]

ax.legend(handles=legend_elements, loc='lower center', ncol=3,

          facecolor=PANEL, edgecolor=BORDER, labelcolor=WHITE,

          fontsize=8, bbox_to_anchor=(0.5, -0.01))

plt.tight_layout()

plt.savefig("outputs-7/fig2_search_tree.png",

           dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig2_search_tree.png")

# =====

# FIGURE 3 - ALPHA-BETA PRUNING DIAGRAM

# =====

def fig_alpha_beta():

    """

    Hand-crafted 3-level alpha-beta example tree showing exact

     $\alpha/\beta$  updates and where pruning fires.

    """

    fig, ax = plt.subplots(figsize=(16, 9), facecolor=BG)

    ax.set_facecolor(BG); ax.axis('off')

    fig.suptitle("ALPHA-BETA PRUNING - Step-by-Step Walkthrough",

                fontsize=16, fontweight='black', color=WHITE, y=0.97)

```

```

# Node definitions: (id, x, y, label, score, is_max, pruned)

nodes = {

    # root

    0: (0.50, 0.82, "MAX", None, True, False),

    # level 1

    1: (0.20, 0.58, "MIN", None, False, False),

    2: (0.50, 0.58, "MIN", None, False, False),

    3: (0.80, 0.58, "MIN", None, False, False),

    # level 2 - leaves

    4: (0.10, 0.30, "", 3, True, False),

    5: (0.20, 0.30, "", 5, True, False),

    6: (0.30, 0.30, "", 2, True, False),

    7: (0.40, 0.30, "", 9, True, False),

    8: (0.50, 0.30, "", 0, True, False),

    9: (0.60, 0.30, "", 7, True, True ), # PRUNED

    10: (0.70, 0.30, "", 5, True, False),

    11: (0.80, 0.30, "", 8, True, True ), # PRUNED

    12: (0.90, 0.30, "", 6, True, True ), # PRUNED

}

# Edges: (parent, child)

edges = [(0,1), (0,2), (0,3),

         (1,4), (1,5), (1,6),

         (2,7), (2,8), (2,9),

         (3,10), (3,11), (3,12)]

```

```

# Computed scores bottom-up

scores = {4:3,5:5,6:2, 7:9,8:0, 10:5}

min_scores = {1: min(3,5,2), 2: min(9,0), 3: 5} # node 9 pruned

scores.update(min_scores)

scores[0] = max(3,0,5) # = 5


# Draw edges

for p,c in edges:

    x0,y0,*_ = nodes[p]; x1,y1,*_ = nodes[c]

    pr = nodes[c][5]

    ax.annotate("", xy=(x1,y1+0.045), xytext=(x0,y0-0.055),

                  arrowprops=dict(arrowstyle="-|>",

                                   color=ORANGE if pr else BORDER,

                                   lw=1.5 if pr else 1.0,

                                   connectionstyle="arc3,rad=0.0",

                                   mutation_scale=10))

    if pr:

        mx = (x0+x1)/2; my = (y0+y1)/2

        ax.text(mx+0.01, my, "X PRUNED",

                fontsize=7, color=ORANGE, fontweight='bold')


# Draw nodes

for nid,(x,y,lbl,leaf_sc,is_max,pruned) in nodes.items():

    sc = leaf_sc if leaf_sc is not None else scores.get(nid)

```

```

ec = ORANGE if pruned else (ACCENT_X if is_max else ACCENT_O)

fc = "#2A1A00" if pruned else ("#1C0505" if is_max else "#05191C")

radius = 0.045

circ = Circle((x,y), radius, facecolor=fc, edgecolor=ec,

               linewidth=2.0, transform=ax.transData, zorder=3)

ax.add_patch(circ)

# role label inside

if lbl:

    ax.text(x, y+0.01, lbl, ha='center', va='center',

            fontsize=9, color=ec, fontweight='black',

            transform=ax.transData, zorder=4)

# score outside node

if sc is not None:

    sc_clr = GREEN if sc>0 else (ACCENT_O if sc<0 else GREY)

    ax.text(x, y - radius - 0.04, str(sc),

            ha='center', va='top', fontsize=11,

            color=sc_clr, fontweight='black',

            transform=ax.transData)

# leaf: show raw value

if leaf_sc is not None and not pruned:

    ax.text(x, y, str(leaf_sc), ha='center', va='center',

            fontsize=10, color=WHITE, fontweight='black',

```

```

transform=ax.transData, zorder=5)

# Alpha-beta annotations

annotations = [

    (0.50, 0.73, "α=3 after left subtree\nβ=+∞", ACCENT_X),

    (0.20, 0.46, "min(3,5,2)=3\n→ parent α=3", ACCENT_O),

    (0.50, 0.46, "min(9,0)=0 < α=3\nNode 9 pruned!", ORANGE),

    (0.80, 0.46, "5 ≥ α=3 → no pruning\n11,12 still pruned (β≤α)", ORANGE),

]

for ax_,ay_,txt,clr in annotations:

    ax_.text(ax_, ay_-0.06, txt, ha='center', va='top',

             fontsize=7.5, color=clr, linespacing=1.4,

             bbox=dict(boxstyle='round,pad=0.3', facecolor=PANEL,

                       edgecolor=clr, linewidth=0.8),

             transform=ax.transAxes)

# Step legend

steps_txt = (

    "STEP-BY-STEP:\n"

    "① Evaluate leftmost subtree (leaves 4,5,6) → MIN node 1 = 3\n"

    "② Root MAX updates α = max(-∞, 3) = 3\n"

    "③ Evaluate subtree of node 2: leaf 7=9, leaf 8=0 → partial min=0\n"

    "    β(node 2) = 0 ≤ α(root) = 3 → PRUNE leaf 9\n"

    "④ Evaluate subtree of node 3: leaf 10=5\n"

    "    β(node 3) = 5 ≥ α=3 → continue (but leaves 11,12 ≤ 5 → pruned)\n"

```

```

    "⑤ Root MAX = max(3, 0, 5) = 5 → choose rightmost branch"

)

ax.text(0.02, 0.15, steps_txt, transform=ax.transAxes,

        fontsize=9, color=WHITE, linespacing=1.6, va='top',

        bbox=dict(boxstyle='round,pad=0.5', facecolor=PANEL,

                  edgecolor=PURPLE, linewidth=1.2))

# Level labels

for y_, txt, clr in [(0.82,"Level 0 - MAX (root)", ACCENT_X),

                    (0.58,"Level 1 - MIN",          ACCENT_O),

                    (0.30,"Level 2 - Leaves",        GOLD)]:

    ax.text(0.01, y_, txt, transform=ax.transAxes,

            fontsize=8, color=clr, va='center', fontweight='bold')

plt.tight_layout()

plt.savefig("outputs-7/fig3_alpha_beta.png",

            dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig3_alpha_beta.png")

# =====

# FIGURE 4 - FULL GAME SIMULATION (AI vs AI)

# =====

```

```

def simulate_game():

    """Simulate an AI-vs-AI game, returns list of (board_snapshot, move, player)."""

    board = [' ']*9

    history = [(board[:], None, None)]

    turn = 0

    while not is_terminal(board):

        player = 'X' if turn % 2 == 0 else 'O'

        is_max = (player == 'X')

        # For O (minimiser) we negate

        if is_max:

            mv, _ = best_move(board)

        else:

            # Best move for O

            best_sc = math.inf; mv = None

            for m in get_available_moves(board):

                b2 = board[:]; b2[m] = 'O'

                s = minimax(b2, 0, True)

                if s < best_sc:

                    best_sc, mv = s, m

            board[mv] = player

        history.append((board[:], mv, player))

        turn += 1

    return history

def fig_game_simulation():

```

```

history = simulate_game()

n = len(history)    # typically 6-9 steps + initial = up to 10 frames

cols = 5

rows = math.ceil(n / cols)

fig = plt.figure(figsize=(cols*3.2, rows*3.5), facecolor=BG)

fig.suptitle("FULL GAME SIMULATION - Minimax X vs Minimax O (Optimal Play)",
             fontsize=14, fontweight='black', color=WHITE, y=0.99)

for i, (board, move, player) in enumerate(history):

    ax = fig.add_subplot(rows, cols, i+1)

    winner, wl = check_winner(board)

    hl = [move] if move is not None else []

    title_clr = GOLD if i==0 else (ACCENT_X if player=='X' else ACCENT_O)

    if winner:

        title_clr = GREEN

        title = ("Initial Board" if i==0

                 else f"Move {i}: {player} → cell {move}"

                 + (" ✓ WINS!" if winner else ""))

    if i == len(history)-1 and not winner:

        title = f"Move {i}: {player} → cell {move} · DRAW ·"

        title_clr = GREY

    draw_board_on_ax(ax, board,

                     highlight_cells=hl,

```



```

        win_line=list(wl) if wl else None,

        title=title, title_color=title_clr,

        fontsize=8)

plt.tight_layout(rect=[0, 0, 1, 0.97])

plt.savefig("outputs-7/fig4_game_simulation.png",

            dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig4_game_simulation.png")

# =====

# FIGURE 5 - HEATMAP: best-move frequency from all opening positions

# =====

def fig_heatmap():

    """Show how often each cell is chosen as the best first move."""

    counts = np.zeros(9)

    # From the empty board, X always picks the same cell, but let's show

    # best moves from all single-O boards (second ply)

    for o_pos in range(9):

        board = [' ']*9

        board[o_pos] = 'O'

        mv, _ = best_move(board)

        counts[mv] += 1

```

```

fig, axes = plt.subplots(1, 2, figsize=(13, 5.5), facecolor=BG)

fig.suptitle("BEST-MOVE HEATMAP - X's Optimal Response to Every O Opening",
             fontsize=13, fontweight='black', color=WHITE, y=1.01)

# Left: heatmap grid

ax = axes[0]; ax.set_facecolor(BG); ax.axis('off')

ax.set_title("Frequency X chooses each cell\n(across all 9 O opening moves)",
            color=GOLD, fontsize=10)

cmap = plt.cm.YlOrRd

norm_counts = counts / counts.max()

for idx in range(9):

    row, col = divmod(idx, 3)

    y = 2 - row; x = col

    clr = cmap(norm_counts[idx])

    rect = FancyBboxPatch((x+0.05, y+0.05), 0.9, 0.9,
                           boxstyle="round,pad=0.05",
                           facecolor=clr, edgecolor=WHITE, linewidth=0.8)

    ax.add_patch(rect)

    ax.text(x+0.5, y+0.55, str(int(counts[idx])),
            ha='center', va='center', fontsize=20,
            color='white' if norm_counts[idx]>0.4 else BG,
            fontweight='black')

    ax.text(x+0.5, y+0.20, f"cell {idx}",
            ha='center', va='center', fontsize=8,

```

```

        color='white' if norm_counts[idx]>0.4 else BG)

ax.set_xlim(0,3); ax.set_ylim(0,3); ax.set_aspect('equal')

# Colourbar

sm = plt.cm.ScalarMappable(cmap=cmap,

                           norm=plt.Normalize(0, counts.max()))

sm.set_array([])

cb = plt.colorbar(sm, ax=ax, fraction=0.03, pad=0.04)

cb.set_label("Times chosen", color=WHITE, fontsize=8)

cb.ax.yaxis.set_tick_params(color=WHITE)

plt.setp(cb.ax.yaxis.get_ticklabels(), color=WHITE)


# Right: bar chart

ax2 = axes[1]; ax2.set_facecolor(PANEL)

cell_labels = [f"Cell {i}" for i in range(9)]

bar_colors = [cmap(norm_counts[i]) for i in range(9)]

bars = ax2.bar(cell_labels, counts, color=bar_colors, edgecolor=BORDER,
linewidth=0.8)

ax2.set_facecolor(PANEL)

ax2.set_title("Bar Chart of Best-Move Frequency", color=GOLD, fontsize=10)

ax2.set_xlabel("Cell Index (0-8)", color=GREY, fontsize=9)

ax2.set_ylabel("Frequency", color=GREY, fontsize=9)

ax2.tick_params(colors=WHITE, rotation=30)

ax2.spines[:].set_color(BORDER)

for spine in ax2.spines.values():

```

```

spine.set_edgecolor(BORDER)

for bar, cnt in zip(bars, counts):

    if cnt:

        ax2.text(bar.get_x()+bar.get_width()/2, bar.get_height()+0.05,

                 str(int(cnt)), ha='center', va='bottom',

                 color=WHITE, fontsize=9, fontweight='bold')

ax2.set_ylim(0, counts.max()*1.25)

ax2.text(0.98, 0.97,

        "Cell 4 (centre) is\nX's most common reply",

        transform=ax2.transAxes, ha='right', va='top',

        fontsize=8.5, color=GOLD,

        bbox=dict(boxstyle='round', facecolor=BG, edgecolor=GOLD, lw=1))

plt.tight_layout()

plt.savefig("outputs-7/fig5_heatmap.png",

           dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig5_heatmap.png")

```

```

# =====
# FIGURE 6 - NODE COUNT: Minimax vs Alpha-Beta comparison
# =====

```

```

_node_counter = [0]

def minimax_count(board, depth, is_max, counter, alpha=-math.inf, beta=math.inf,
ab=True):

    counter[0] += 1

    winner, _ = check_winner(board)

    if winner == 'X': return 10-depth

    if winner == 'O': return depth-10

    if ' ' not in board: return 0

    moves = get_available_moves(board)

    best = -math.inf if is_max else math.inf

    for mv in moves:

        b2 = board[:]; b2[mv] = 'X' if is_max else 'O'

        s = minimax_count(b2, depth+1, not is_max, counter, alpha, beta, ab)

        if is_max:

            best = max(best,s)

            if ab: alpha = max(alpha,s)

        else:

            best = min(best,s)

            if ab: beta = min(beta,s)

            if ab and beta<=alpha: break

    return best

def fig_node_comparison():

```

```

boards = [

    [' ']*9,

    ['X',' ',' ',' ',' ',' ',' ',' ',' ',' '],

    ['X','O',' ',' ',' ','X',' ',' ',' ',' '],

    ['X','O','X','O','X',' ',' ',' ',' ','O'],

]

labels = ["Empty\n(0 moves)", "1 move\nplayed", "4 moves\nplayed", "6
moves\nplayed"]

plain_counts = []

ab_counts = []

for b in boards:

    c1=[0]; c2=[0]

    minimax_count(b,0,True,c1,ab=False)

    minimax_count(b,0,True,c2,ab=True)

    plain_counts.append(c1[0])

    ab_counts.append(c2[0])

x = np.arange(len(labels))

w = 0.35

fig, axes = plt.subplots(1, 2, figsize=(13, 5.5), facecolor=BG)

fig.suptitle("NODES EVALUATED: Plain Minimax vs Alpha-Beta Pruning",

             fontsize=13, fontweight='black', color=WHITE, y=1.01)

```

```

# Bar chart

ax = axes[0]; ax.set_facecolor(PANEL)

ax.bar(x-w/2, plain_counts, width=w, color=ACCENT_X, edgecolor=BG,

       label="Plain Minimax")

ax.bar(x+w/2, ab_counts, width=w, color=ACCENT_O, edgecolor=BG,

       label="Alpha-Beta")

for i,(p,a) in enumerate(zip(plain_counts,ab_counts)):

    ax.text(i-w/2, p+100, f"{p:,.}", ha='center', fontsize=7.5,

            color=ACCENT_X, fontweight='bold')

    ax.text(i+w/2, a+100, f"{a:,.}", ha='center', fontsize=7.5,

            color=ACCENT_O, fontweight='bold')

ax.set_xticks(x); ax.set_xticklabels(labels, color=WHITE, fontsize=9)

ax.set_ylabel("Nodes evaluated", color=GREY, fontsize=9)

ax.set_title("Node Count Comparison", color=GOLD, fontsize=10)

ax.legend(facecolor=BG, labelcolor=WHITE, edgecolor=BORDER)

ax.tick_params(colors=WHITE)

ax.spines[:].set_color(BORDER)

ax.set_facecolor(PANEL)


# Reduction % line chart

ax2 = axes[1]; ax2.set_facecolor(PANEL)

reduction = [(1 - a/p)*100 if p else 0 for p,a in zip(plain_counts, ab_counts)]

ax2.plot(labels, reduction, color=GOLD, lw=2.5, marker='o', markersize=9,

         markerfacecolor=BG, markeredgecolor=GOLD, markeredgewidth=2)

for i,(lbl,r) in enumerate(zip(labels, reduction)):

```

```

        ax2.text(i, r+1.5, f"{r:.0f}%", ha='center', fontsize=9,
                 color=GOLD, fontweight='bold')

ax2.set_ylim(0, 100); ax2.set_xlabel("Board State", color=GREY, fontsize=9)

ax2.set_ylabel("Nodes Pruned (%)", color=GREY, fontsize=9)

ax2.set_title("Pruning Efficiency (%)", color=GOLD, fontsize=10)

ax2.tick_params(colors=WHITE)

ax2.spines[:].set_color(BORDER)

ax2.axhline(50, color=BORDER, lw=0.8, linestyle='--')

ax2.text(3.05, 50, "50%", color=GREY, fontsize=8, va='center')


plt.tight_layout()

plt.savefig("outputs-7/fig6_node_comparison.png",
           dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig6_node_comparison.png")


# =====
# FIGURE 7 - SCORE LANDSCAPE (all possible first moves)
# =====


def fig_score_landscape():

    fig, axes = plt.subplots(1, 3, figsize=(14, 5), facecolor=BG)

    fig.suptitle("MINIMAX SCORE LANDSCAPE - X's Score for Every Possible Move",
                fontsize=13, fontweight='black', color=WHITE, y=1.01)

```



```

scenarios = [

    ([' ']*9, True, "Empty Board - X first"),

    (['O',' ',' ',' ',' ',' ',' ',' ',' '], True, "After O→0"),

    (['O',' ',' ',' ',' ','X',' ',' ',' ','O'], True, "After O→0,X→4,O→8"),

]

```

```

for ax, (board, is_max, title) in zip(axes, scenarios):

    ax.set_facecolor(BG); ax.axis('off')

    ax.set_title(title, color=GOLD, fontsize=9, fontweight='bold')

    scores_grid = np.full(9, np.nan)

    for mv in get_available_moves(board):

        b2 = board[:]; b2[mv] = 'X' if is_max else 'O'

        s = minimax(b2, 0, not is_max)

        scores_grid[mv] = s

    cmap = plt.cm.RdYlGn

    finite = scores_grid[~np.isnan(scores_grid)]

    vmin = finite.min() if len(finite) else -10

    vmax = finite.max() if len(finite) else 10

    for idx in range(9):

        row, col = divmod(idx, 3)

        y = 2-row; x = col

        val = scores_grid[idx]

```

```

if np.isnan(val):

    # occupied cell

    fc = "#0A0A0A"; txt = board[idx]

    txt_clr = ACCENT_X if txt=='X' else ACCENT_O

    fs = 18

else:

    norm_v = (val-vmin)/(vmax-vmin+1e-9)

    fc = cmap(norm_v)

    txt = f"{val:+.0f}" if val != 0 else "0"

    txt_clr = BG

    fs = 14

rect = FancyBboxPatch((x+0.05, y+0.05), 0.9, 0.9,

                       boxstyle="round,pad=0.05",

                       facecolor=fc, edgecolor=WHITE, linewidth=0.6)

ax.add_patch(rect)

ax.text(x+0.5, y+0.5, txt,

        ha='center', va='center', fontsize=fs,

        color=txt_clr, fontweight='black')

# best move star

if not np.isnan(val) and len(finite) and val == finite.max():

    ax.text(x+0.82, y+0.82, "★",

            ha='center', va='center', fontsize=11, color=GOLD)

```

```

ax.set_xlim(0,3); ax.set_ylim(0,3); ax.set_aspect('equal')

# Shared colorbar

sm2 = plt.cm.ScalarMappable(cmap=plt.cm.RdYlGn,
                             norm=plt.Normalize(-10,10))

sm2.set_array([])

cb = plt.colorbar(sm2, ax=axes, fraction=0.015, pad=0.03)

cb.set_label("Minimax Score", color=WHITE, fontsize=9)

cb.ax.yaxis.set_tick_params(color=WHITE)

plt.setp(cb.ax.yaxis.get_ticklabels(), color=WHITE)

plt.tight_layout()

plt.savefig("outputs-7/fig7_score_landscape.png",
            dpi=150, bbox_inches='tight', facecolor=BG)

plt.close()

print("✓ fig7_score_landscape.png")

# =====
# MAIN
# =====

if __name__ == "__main__":

    print("Generating all figures...")

    fig_overview()

```

```
fig_search_tree()

fig_alpha_beta()

fig_game_simulation()

fig_heatmap()

fig_node_comparison()

fig_score_landscape()

print("\n✅ All 7 figures saved to outputs-7/")
```

Output:

### MINIMAX ALGORITHM – Tic Tac Toe

#### Pseudocode

```
function MINIMAX(state, depth, isMaximising):
    if TERMINAL(state):
        return SCORE(state, depth)
    if isMaximising:
        best ← -∞
        for move in MOVES(state):
            val ← MINIMAX(APPLY(state, move), depth+1, FALSE)
            best ← MAX(best, val)
        return best
    else: # minimising
        best ← ∞
        for move in MOVES(state):
            val ← MINIMAX(APPLY(state, move), depth+1, TRUE)
            best ← MIN(best, val)
        return best
```

#### Scoring Function

|         |             |
|---------|-------------|
| X wins  | +10 - depth |
| O wins  | depth - 10  |
| Draw    | 0           |
| X plays | Maximise +  |
| O plays | Minimise -  |

#### Alpha-Beta Pruning

Improvement over plain Minimax.

$\alpha$  = best score Maximiser can guarantee  
 $\beta$  = best score Minimiser can guarantee

Prune when  $\beta \leq \alpha$   
(remaining branches can't affect result)

Complexity:  
Plain Minimax:  $O(b^m)$   
Alpha-Beta:  $O(b^{\lceil m/2 \rceil})$  (best case)

For Tic-Tac-Toe:  
 $b = 5, m = 9 \rightarrow 255,168$  positions  
With pruning: significantly fewer nodes

#### Game Tree Facts

|                  |             |
|------------------|-------------|
| Board states     | ~5,478      |
| Terminal nodes   | 255,168     |
| X wins           | 131,184     |
| O wins           | 77,904      |
| Draws            | 46,080      |
| Tree depth       | max 9       |
| Branching factor | $\leq 9$    |
| Result (optimal) | Always draw |

Example: X to move → best=10

|   |   |   |
|---|---|---|
| X |   | O |
|   | X |   |
| O |   |   |

After best move → X wins!

|   |   |   |
|---|---|---|
| X |   | O |
|   | X |   |
| O |   | X |

■ Maximiser (X)

■ Minimiser (O)





